

ST-540 Assignment-4

Tilekbek Zhoroev

Problem 1 Give an advantage and a disadvantage of the following methods:

- a. Maximum a posterior estimation: *Fast but doesn't quantify uncertainty*
- b. Numerical integration: *Accurate but not feasible in high dimensions*
- c. Bayesian central limit theorem: *Fast but need large sample size.*
- d. Gibbs sampling: *Works in high dimensions and can be accurate but requires known conjugate priors.*
- e. Metropolis Hastings sampling: *Flexible but requires tuning.*

Problem 2

Consider the model $Y_i|\mu, \sigma^2 \sim \text{Normal}(\mu, \sigma^2)$ for $i = 1, \dots, n$ and $Y_i|\mu, \delta, \sigma^2 \sim \text{Normal}(\mu + \delta, \sigma^2)$ for $i = n+1, \dots, n+m$, where $\mu, \delta \sim \text{Normal}(\mu, 100^2)$ and $\sigma^2 \sim \text{InvGamma}(0.01, 0.01)$

- a. Give an example of a real experiment for which this would be an appropriate model.

Y_i are the average monthly amount of money one person consumes for daily expenses. After COVID-19, we would like to analyse the effect of pandemic in m month to person's daily expenses.

Another example: Weights of two groups of people, n people with age 10-20 and m people with age 20-30

- b. Derive the full conditional posterior distributions for μ, δ and σ^2

First, by given expressions we can obtain that

$$\begin{aligned}
f(Y_1, Y_2, \dots, Y_{m+n}|\mu, \delta, \sigma^2) &= \prod_{i=1}^n f(Y_i|\mu, \sigma^2) \prod_{i=n+1}^{n+m} f(Y_i|\mu + \delta, \sigma^2) \\
&= \prod_{i=1}^n \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(Y_i - \mu)^2}{2\sigma^2}\right) \prod_{i=n+1}^{n+m} \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(Y_i - \mu - \delta)^2}{2\sigma^2}\right) \\
&\propto \exp\left(-\frac{1}{2\sigma^2} \sum_{i=1}^n (Y_i - \mu)^2\right) \times \exp\left(-\frac{1}{2\sigma^2} \sum_{i=n+1}^{n+m} (Y_i - \mu - \delta)^2\right) \\
&\propto \exp\left(-\frac{(m+n)\mu - 2\mu(\sum_{i=1}^n Y_i + \sum_{i=1}^n (Y_i - \delta))}{2\sigma^2}\right)
\end{aligned}$$

Then, from section 2.1.3 we have,

$$\begin{aligned}
f(\mu|Y_1, \dots, Y_{n+m}, \delta, \sigma^2) &= \frac{f(Y_1, \dots, Y_{n+m}, \mu, \delta, \sigma^2)}{f(Y_1, \dots, Y_{n+m}, \mu, \delta, \sigma^2)} \\
&\propto f(Y_1, \dots, Y_{n+m}, \mu, \delta, \sigma^2) \\
&\propto f(Y_1, \dots, Y_{n+m}|\mu, \delta, \sigma^2)f(\mu) \\
&\propto \text{Normal} \left(\frac{\sum_{i=1}^n Y_i + \sum_{i=n+1}^{m+n} (Y_i - \delta)}{n + m + \sigma^2/100^2}, \frac{1}{(n + m)/\sigma^2 + 1/100^2} \right).
\end{aligned}$$

Since δ involved only after $i = n + 1$,

$$\delta|Y, \mu, \sigma^2 \propto \text{Normal} \left(\frac{\sum_{i=n+1}^{m+n} (Y_i - \mu)}{m + \sigma^2/100^2}, \frac{1}{m/\sigma^2 + 1/100^2} \right).$$

To find the full condition distribution for σ^2 we would use the fact from section 2.1.4 and we get

$$\sigma^2|Y, \mu, \delta \propto \text{InvGamma} \left(0.01 + (n + m)/2, \sum_{i=1}^n (Y_i - \mu)^2/2 + \sum_{i=n+1}^{n+m} (Y_i - \mu - \delta)^2/2 + 0.01 \right)$$

- c. Simulate a dataset from this model with $n = m = 50$, $\mu = 10$, $\delta = 1$, and $\sigma = 2$. Write your own Gibbs sampling code (not in JAGS) to fit the model above to the simulated data and plot the marginal posterior density for each parameter. Are you able to recover the true values reasonably well?

```

set.seed(42)
# Given parameters
n <- 50; m <- 50; mu_0 <- 10; delta_0 <- 1; sigma_0 <- 2
#generate likelihoods
Y1 <- rnorm(n,mu_0,sigma_0); Y2 <- rnorm(m,mu_0+delta_0,sigma_0); Y <- c(Y1,Y2)

# Prep for Gibbs sampling
S       <- 10^5 # number of iterations
#priors
pri_var <- 100^2;   a <- b <- 0.01; mu <- mean(Y); delta <- mean(Y)-mean(Y2); sigma2 <- var(Y)

Gibss<-matrix(0,S,3)
for(iter in 1:S){
  # mu
  prec <- (n+m)/sigma2 + 1/pri_var
  mn   <- sum(Y1)/sigma2 + sum(Y2-delta)/sigma2
  mu   <- rnorm(1,mn/prec,1/sqrt(prec))

  # delta
  prec <- m/sigma2 + 1/pri_var
  mn   <- sum(Y2-mu)/sigma2
  delta <- rnorm(1,mn/prec,1/sqrt(prec))

  # sigma2
}
```

```

sigma2 <- rinvgamma(1,a+(n+m)/2,(sum((Y1-mu)^2) + sum((Y2-mu-delta)^2))/2+b)

Gibss[iter,] <- c(mu,delta,sigma2)
}
Gibss<-data.frame(Gibss)
colnames(Gibss) = c("mu","delta","sigma^2")

Gibss%>%
  ggplot(aes(x=seq(1,S), y = mu ))+geom_line()+xlab("iteration")

Gibss%>%
  ggplot(aes(mu))+  

  geom_histogram(aes(y = ..density..), colour = "white", bins = 40) +  

  geom_density(colour ="blue")+
  geom_vline(xintercept=mu_0, linetype="dashed", colour = "red")

Gibss%>%
  ggplot(aes(x=seq(1,S), y = delta ))+geom_line()+xlab("iteration")

Gibss%>%
  ggplot(aes(delta))+  

  geom_histogram(aes(y = ..density..), color = "white", bins = 40) +  

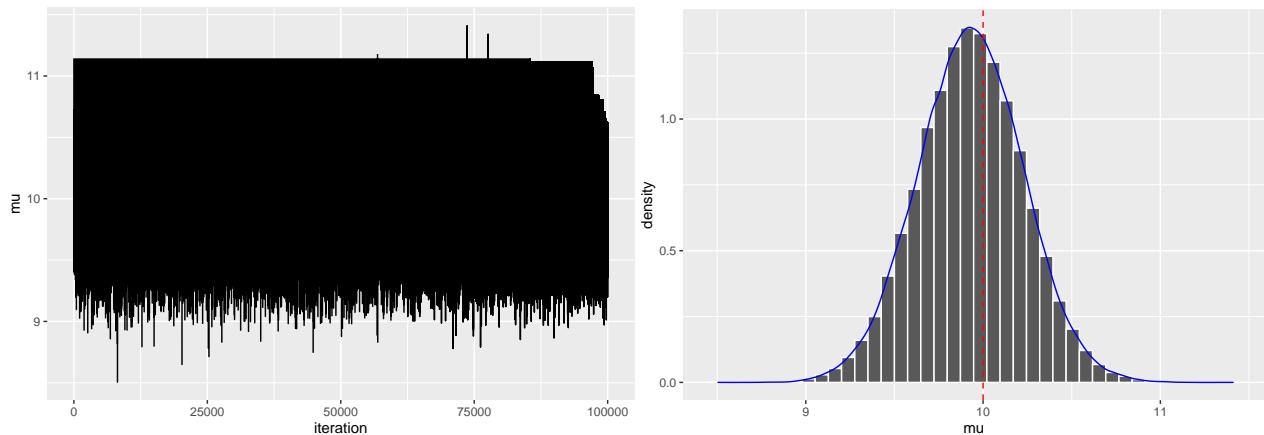
  geom_density(colour ="blue")+
  geom_vline(xintercept=delta_0, linetype="dashed", color = "red")

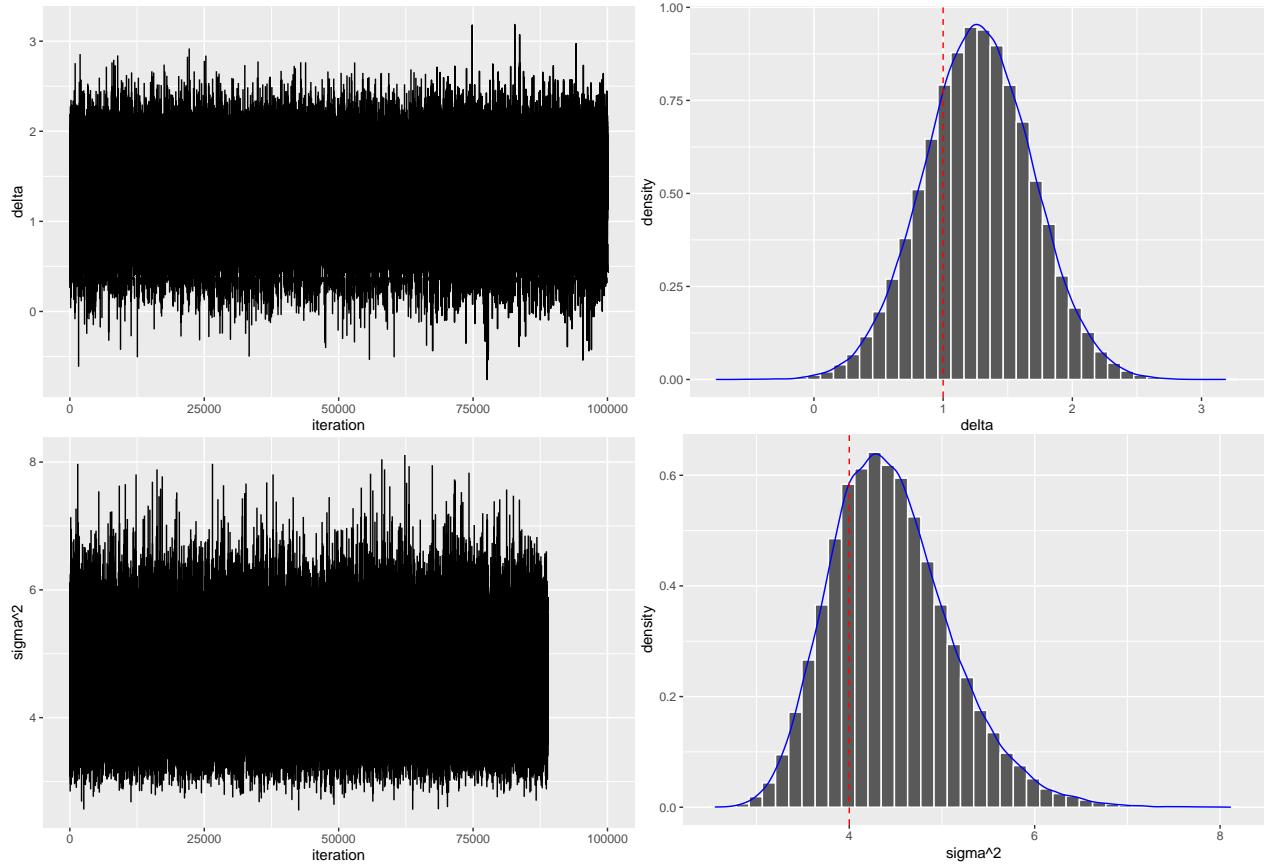
Gibss%>%
  ggplot(aes(x=seq(1,S), y = `sigma^2` ))+geom_line()+xlab("iteration")

Gibss%>%
  ggplot(aes(`sigma^2`)) +
  geom_histogram(aes(y = ..density..), color = "white", bins = 40) +  

  geom_density(colour ="blue") +
  geom_vline(xintercept=sigma_0^2, linetype="dashed", color = "red")

```





Problem 3

Fit the following model to the NBA free throw data in the table in the Exercise 17 in Chapter 1:

Player	Overall proportion	Clutch makes	Clutch attempts
Russell Westbrook	0.845	64	75
James Harden	0.847	72	95
Kawhi Leonard	0.880	55	63
LeBron James	0.674	27	39
Isaiah Thomas	0.909	75	83
Stephen Curry	0.898	24	26
Giannis Antetokounmpo	0.770	28	41
John Wall	0.801	66	82
Anthony Davis	0.802	40	54
Kevin Durant	0.875	13	16

$Y_i|\theta_i \sim \text{Binomial}(n_i; \theta_i)$ and $\theta_i|m \sim \text{Beta}[\exp(m)q_i, \exp(m)(1 - q_i)]$, where Y_i is the number of made clutch shots for player $i = 1, \dots, 10$, n_i is the number of attempted clutch shots, $q_i \in (0, 1)$ is the overall proportion, and $m \sim \text{Normal}(0, 10)$.

- a. Explain why this is a reasonable prior for θ_i .

Since the domain of the beta distribution covers 0 to 1 and the mean of the beta distribution is

$$\frac{e^m q_i}{e^m q_i + e^m (1 - q_i)} = q_i$$

so distribution is centered at q_i . and takes values on $[0, 1]$.

- b. Explain the role of m in the prior.

This determines the spread of the distribution around q_i .

- c. Derive the full conditional posterior for θ_1 .

Since the number of the clutches of each player is independent we have

$$\begin{aligned} f(Y_1, \dots, Y_{10}, \theta_1, \dots, \theta_{10}|m) &= \prod_{i=1}^{10} f(Y_i|\theta_i)f(\theta_i|m) \\ &\propto \prod_{i=1}^{10} \text{Beta}(Y_i + \exp(m)q_i, n_i - Y_i + \exp(m)(1 - q_i)) \end{aligned}$$

where we employ the Beta-Binomial conjugacy. Next, the full conditional posterior of the θ_1 is

$$\begin{aligned} f(\theta_1|Y_1, \dots, Y_{10}, \theta_2, \dots, \theta_{10}, m) &= \frac{f(Y_1, \dots, Y_{10}, \theta_1, \dots, \theta_{10}|m)}{f(Y_1, \dots, Y_{10}, \theta_2, \dots, \theta_{10}|m)} \\ &\propto \text{Beta}(Y_1 + \exp(m)q_1, n_1 - Y_1 + \exp(m)(1 - q_1)) \end{aligned}$$

- d. Write your own MCMC algorithm to compute a table of posterior means and 95% credible intervals for all 11 model parameters $(\theta_1, \dots, \theta_{10}, m)$. Turn in commented code.

Similarly as previous part we obtain the full conditional posterior of each θ_i ,

$$\theta_i|\text{rest} \propto \text{Beta}(Y_i + \exp(m)q_i, n_i - Y_i + \exp(m)(1 - q_i)) \quad i = 1, \dots, 10.$$

However for m we don't have nice form of the full conditional distribution so we combine Gibbs and Metropolis algorithm

```
set.seed(420)
#given values in table
Ys <- c(64, 72, 55, 27, 75, 24, 28, 66, 40, 13)
ns <- c(75, 95, 63, 39, 83, 26, 41, 82, 54, 16)
qs <- c(0.845, 0.847, 0.880, 0.674, 0.909, 0.898, 0.770, 0.801, 0.802, 0.875)
#parameters to start MCMC
m <- 0;
thetas <- qs;
S <- 3*10^4;
burn_in = 10^4;
MCMC<-matrix(0,S-burn_in,11); # save only after burn-in
can_sd <- 0.3
#log posterior for m
log_post_m <- function(theta, Ys, ns, qs, m){
  like <- 0
  for(i in 1:10){
    like = like + dbeta(theta[i], Ys[i] + exp(m)*qs[i], ns[i] - Ys[i] + exp(m)*(1-qs[i]), log = TRUE )
  }
  prior = dnorm(m, 0, 10, log=TRUE)
  return(like + prior)}
for(iter in 1:S){
  # Gibbs for each theta
  for(i in 1:10){
    alpha = Ys[i] + exp(m) * qs[i]
    beta = ns[i] - Ys[i] + exp(m) * (1 - qs[i])
```

```

        thetas[i] = rbeta(1, alpha, beta)
    }
# Metropolis for m
can <- rnorm(1, m, can_sd) #proposal distribution
logR <- log_post_m(thetas, Ys, ns, qs, can) - log_post_m(thetas, Ys, ns, qs, m)
R<-exp(logR)
if(runif(1) < R){
    m <- can
}
if(iter>burn_in){
    MCMC[iter - burn_in,] = c(thetas,m)
}

MCMC<-data.frame(MCMC)

colnames(MCMC)<-c("Russell Westbrook", "James Harden", "Kawhi Leonard", "LeBron James",
"Isaiah Thomas", "Stephen Curry", "Giannis Antetokounmpo", "John Wall",
"Anthony Davis", "Kevin Durant", "m")

```

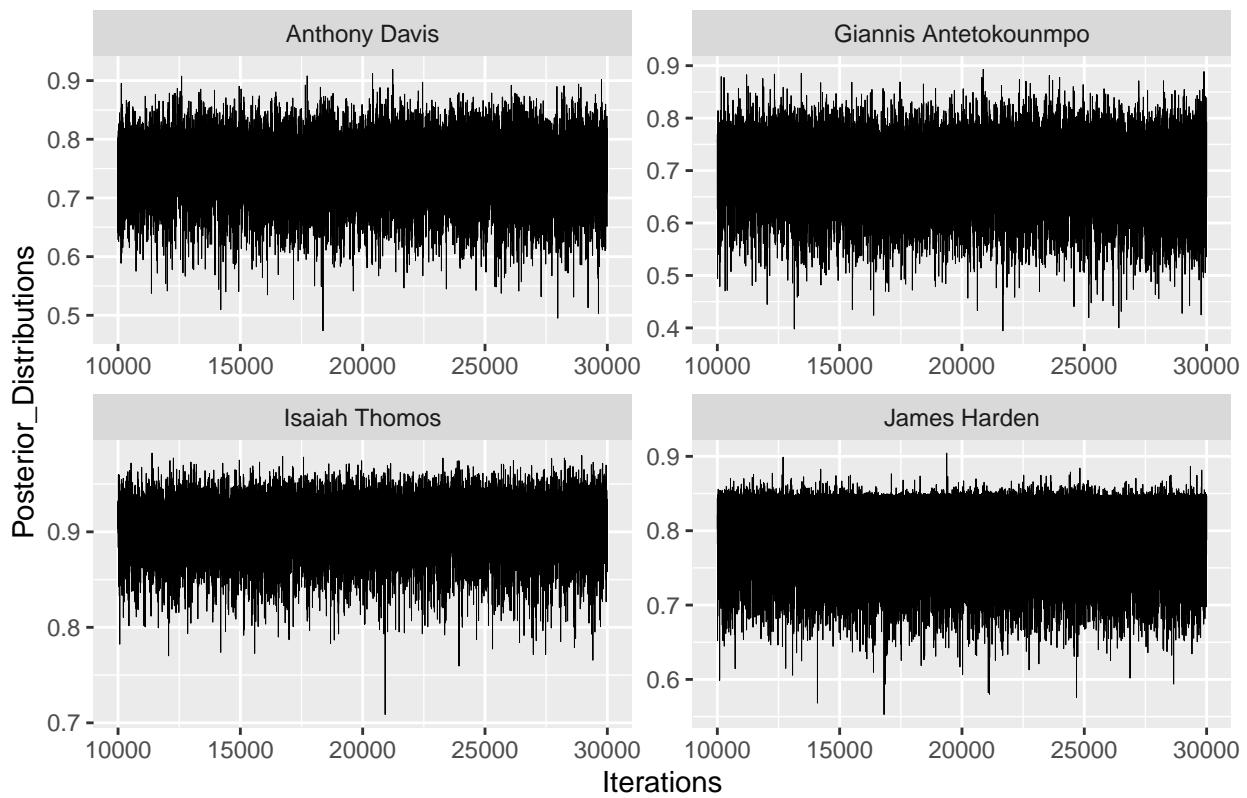
The trace plots of MCMC sample of parameters are presented below.

```

MCMC%>%
pivot_longer(cols = "Russell Westbrook": "m", names_to = "Parameters",
             values_to = "Posterior_Distributions")%>%
ggplot(aes(x=rep(seq(burn_in+1,S), 11), y = Posterior_Distributions))+
xlab("Iterations")+
geom_line(size=.1)+theme(plot.title = element_text(hjust = 0.5))+
ggtitle(" Trace plot of the MCMC samples of each posteriors")+
facet_wrap_paginate(~Parameters,scales = "free", ncol = 2, nrow = 2, page = 1)

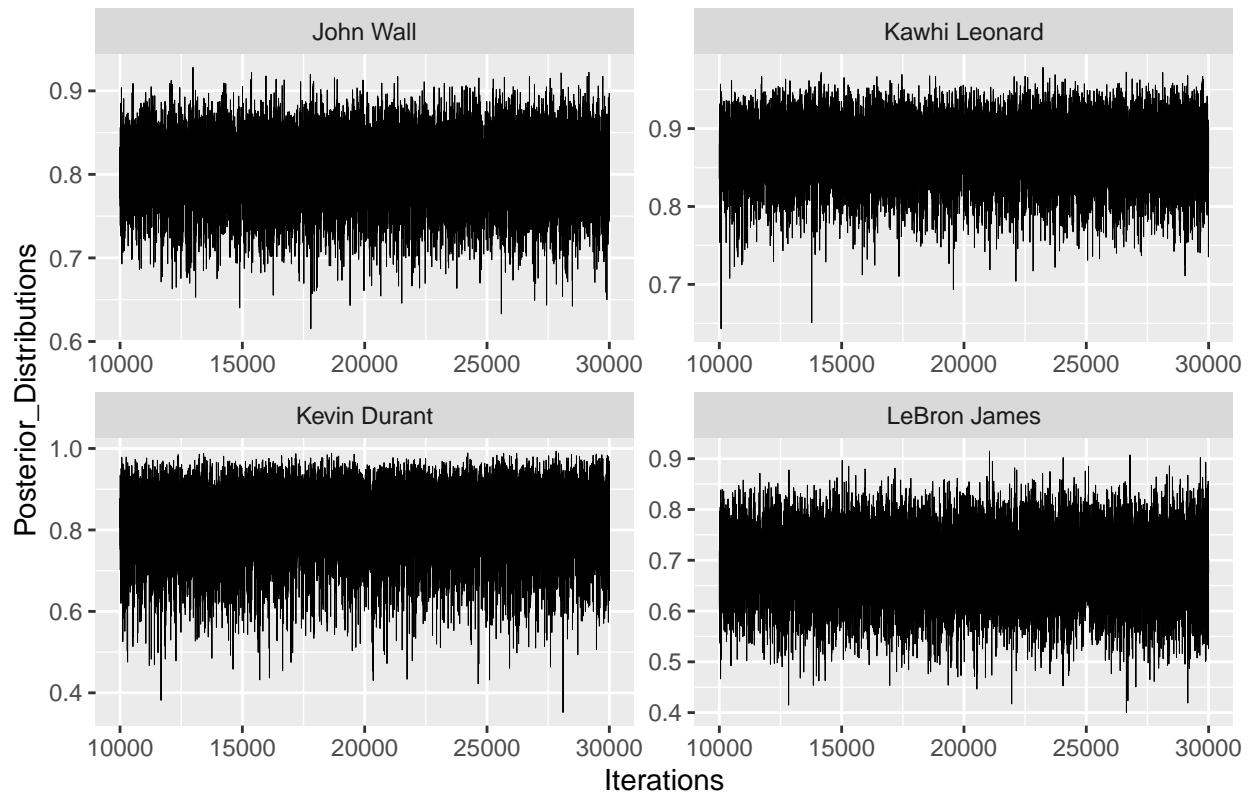
```

Trace plot of the MCMC samples of each posteriors



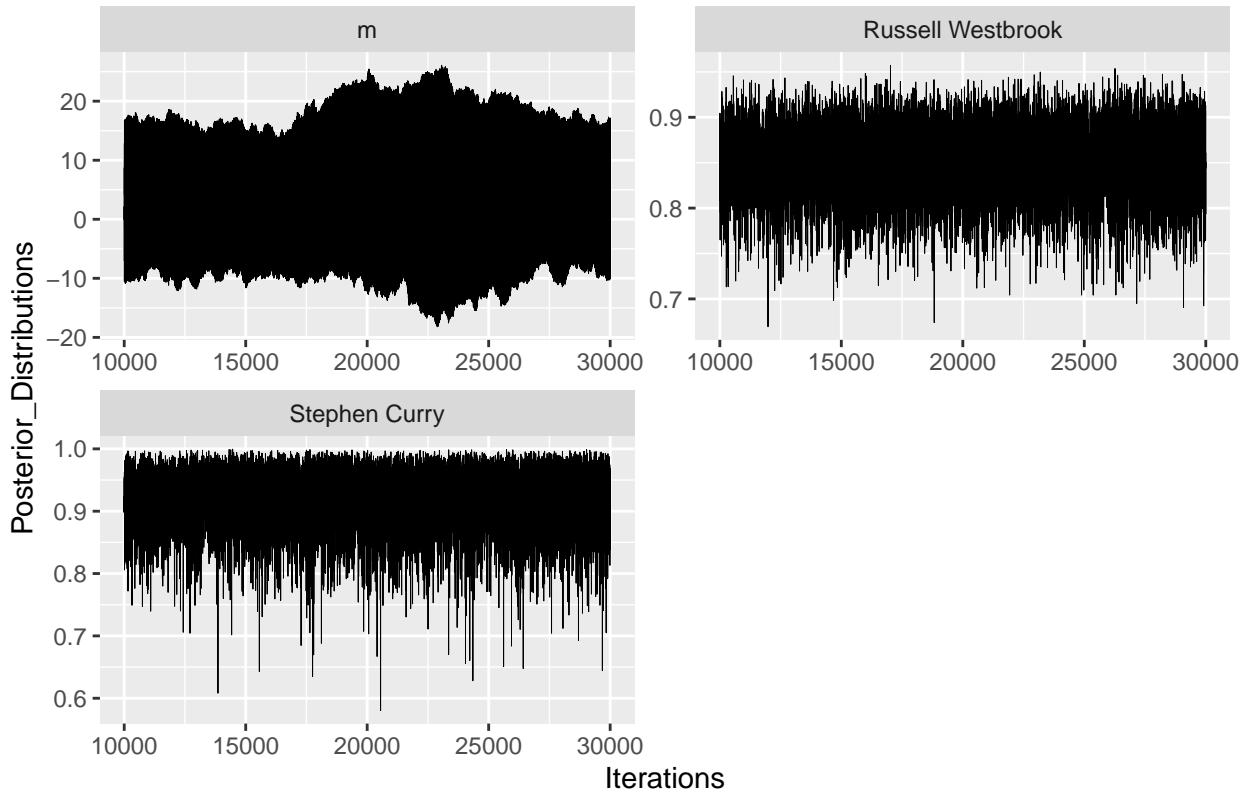
```
MCMC%>%
pivot_longer(cols = "Russell Westbrook": "m", names_to = "Parameters",
             values_to = "Posterior_Distributions")%>%
ggplot(aes(x=rep(seq(burn_in+1,S), 11), y = Posterior_Distributions))+
  xlab("Iterations")+
  geom_line(size=.1)+theme(plot.title = element_text(hjust = 0.5))+  
  ggtitle(" Trace plot of the MCMC samples of each posteriors")+
  facet_wrap_paginate(~Parameters,scales = "free", ncol = 2, nrow = 2, page = 2)
```

Trace plot of the MCMC samples of each posteriors



```
MCMC%>%
pivot_longer(cols = "Russell Westbrook": "m", names_to = "Parameters",
             values_to = "Posterior_Distributions")%>%
ggplot(aes(x=rep(seq(burn_in+1,S), 11), y = Posterior_Distributions))+
  xlab("Iterations")+
  geom_line(size=.09)+theme(plot.title = element_text(hjust = 0.5))+
  ggtitle(" Trace plot of the MCMC samples of each posteriors")+
  facet_wrap_paginate(~Parameters,scales = "free", ncol = 2, nrow = 2, page = 3)
```

Trace plot of the MCMC samples of each posteriors



```
table=sapply(MCMC, quantile, probs = c(.5, 0.025, 0.975))
rownames(table) = c("Means","2.5 %","97.5 %")
knitr::kable(t(table))
```

	Means	2.5 %	97.5 %
Russell Westbrook	0.8451769	0.7782915	0.9143715
James Harden	0.8016995	0.6826536	0.8507454
Kawhi Leonard	0.8799938	0.7967172	0.9347815
LeBron James	0.6742497	0.5654131	0.8053789
Isaiah Thomas	0.9089932	0.8430956	0.9509257
Stephen Curry	0.9001120	0.8230233	0.9853191
Giannis Antetokounmpo	0.7482889	0.5582152	0.8054475
John Wall	0.8010575	0.7276388	0.8738146
Anthony Davis	0.7885786	0.6405154	0.8385163
Kevin Durant	0.8736614	0.6461045	0.9452855
m	2.2746495	-11.5461571	21.4155681

- e. Fit the same model in JAGS. Turn in commented code, and comment on whether the two algorithms returned the same results.

```
#given data
Ys <- c(64,72,55,27,75,24,28,66,40,13)
Ns <- c(75,95,63,39,83,26,41,82,54,16)
qs <- c(0.845, 0.847, 0.880, 0.674, 0.909, 0.898, 0.770, 0.801, 0.802, 0.875)
n = 10
# define string model
model_string <- textConnection("model{
```

```

# Likelihood
for(i in 1:n){
  Y[i] ~ dbin(theta[i], N[i])
}
# Priors
for(i in 1:n){
  theta[i] ~ dbeta(exp(m)*qs[i], exp(m)*(1-qs[i]))
}

m ~ dnorm(0, 10)
}")

# Load the data and compile the MCMC code
inits <- list(theta=qs, m = 0)
data <- list(Y = Ys,N = Ns, qs = qs, n = n)
model <- jags.model(model_string,data = data, inits=inits, n.chains=2)

## Compiling model graph
## Resolving undeclared variables
## Allocating nodes
## Graph information:
##   Observed stochastic nodes: 10
##   Unobserved stochastic nodes: 11
##   Total graph size: 76
##
## Initializing model
#Burn-in for 10000 samples
update(model, 10000, progress.bar="none")
# Generate 20000 post-burn-in samples

params <- c("theta","m")
samples <- coda.samples(model,
                        variable.names=params,
                        n.iter=20000, progress.bar="none")

summary(samples)

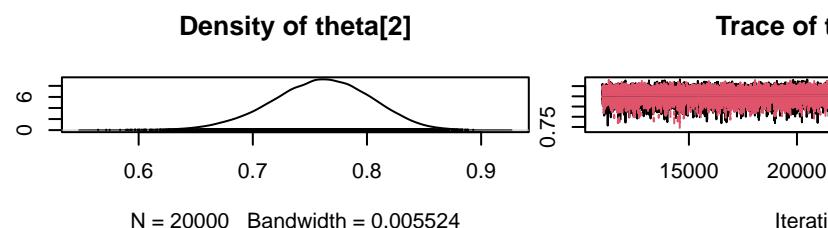
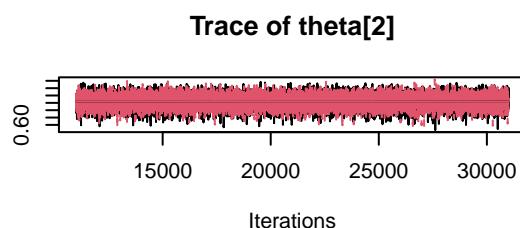
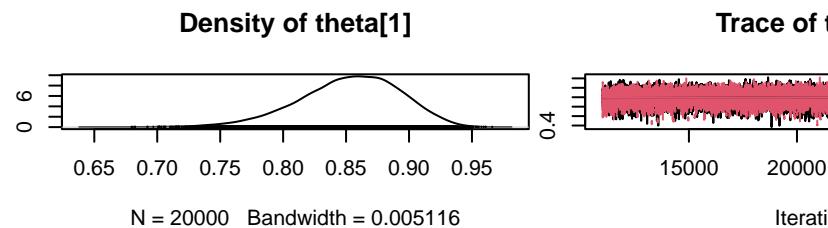
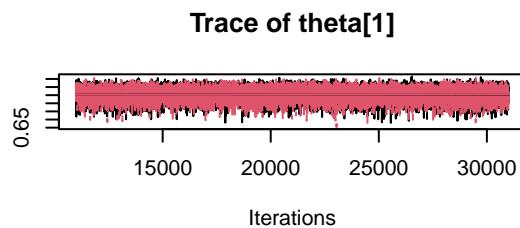
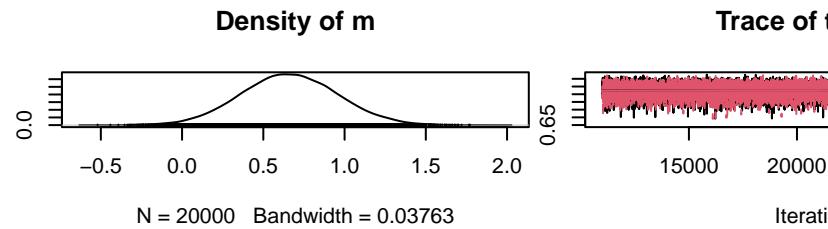
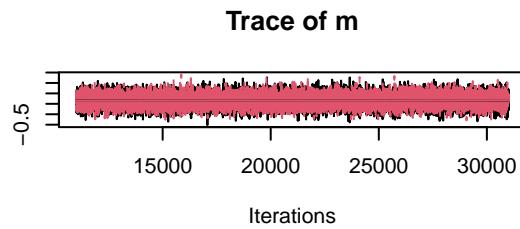
##
## Iterations = 11001:31000
## Thinning interval = 1
## Number of chains = 2
## Sample size per chain = 20000
##
## 1. Empirical mean and standard deviation for each variable,
##    plus standard error of the mean:
##
##           Mean        SD  Naive SE Time-series SE
## m       0.6619  0.29559  0.0014779      0.0018409
## theta[1] 0.8528  0.04018  0.0002009      0.0002670
## theta[2] 0.7596  0.04339  0.0002169      0.0002802
## theta[3] 0.8732  0.04088  0.0002044      0.0002797
## theta[4] 0.6912  0.07113  0.0003557      0.0004594
## theta[5] 0.9039  0.03166  0.0001583      0.0002165
## theta[6] 0.9212  0.05005  0.0002503      0.0004110
## theta[7] 0.6871  0.06953  0.0003477      0.0004457

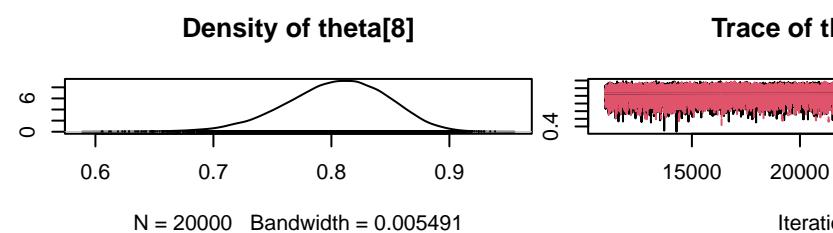
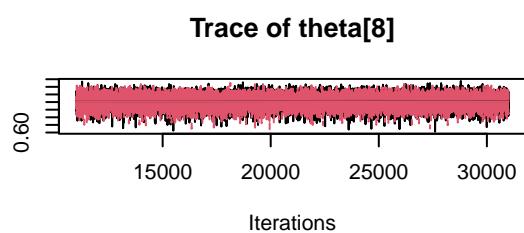
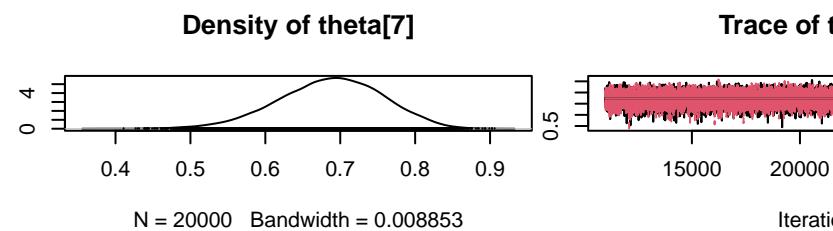
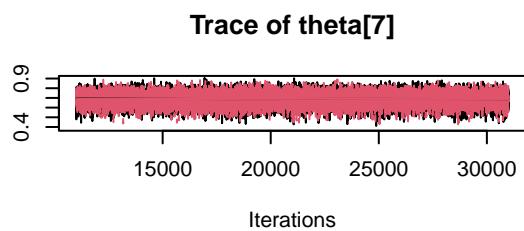
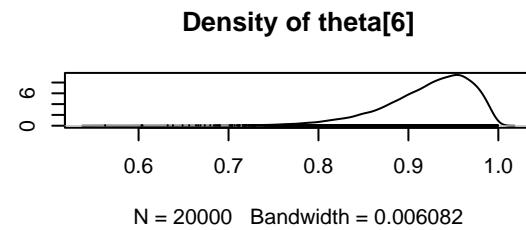
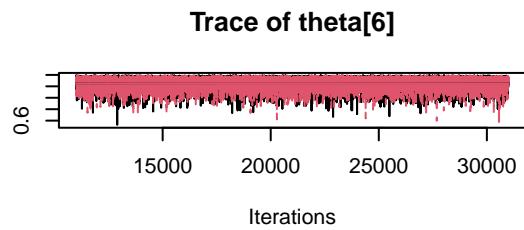
```

```

## theta[8]  0.8043 0.04312 0.0002156      0.0002834
## theta[9]  0.7437 0.05739 0.0002869      0.0003600
## theta[10] 0.8193 0.08812 0.0004406      0.0006229
##
## 2. Quantiles for each variable:
##
##          2.5%   25%   50%   75% 97.5%
## m       0.08821 0.4607 0.6592 0.8624 1.2487
## theta[1] 0.76598 0.8273 0.8557 0.8816 0.9225
## theta[2] 0.67031 0.7312 0.7610 0.7897 0.8399
## theta[3] 0.78215 0.8480 0.8770 0.9026 0.9416
## theta[4] 0.54436 0.6447 0.6943 0.7416 0.8210
## theta[5] 0.83364 0.8843 0.9070 0.9269 0.9564
## theta[6] 0.79901 0.8945 0.9310 0.9585 0.9882
## theta[7] 0.54301 0.6413 0.6899 0.7360 0.8144
## theta[8] 0.71352 0.7767 0.8068 0.8350 0.8809
## theta[9] 0.62524 0.7063 0.7465 0.7845 0.8471
## theta[10] 0.61713 0.7659 0.8312 0.8856 0.9541
plot(samples)

```





- f. What are the advantages and disadvantages of writing your own code as opposed to using JAGS in this problem and in general?

In general using build-in codes are safest since it considered all optimization aspects in terms of time and communication algorithm. Moreover, I may do type when while writing code.